

AN INTRODUCTION TO AKKA PERSISTENCE

Bartosz Mikulski

THE SPEAKER

@mikulskibartosz

- Java / Scala / sometimes JS ;)
- programmer for 5 years
- co-founder of Poznan Scala User Group
meetup.com/poznanscala

EVENT SOURCING

1. An actor receives a command
2. The actor checks whether the command can be applied to the current state
3. The actor generates and stores an event

An event describes the effect of a command

EVENT SOURCING

- commands
- events
- Command Query Responsibility Segregation

EVENT SOURCING - EXTERNAL UPDATES

- if the application replays events, external system processes duplicate messages
- communicate with external system using a Gateway
- disable the Gateway during recovery

EVENT SOURCING - EXTERNAL QUERIES

- returned value influences event processing
- the value may be different during recovery
- use a Gateway to query external systems
- the Gateway may store value returned from the system

PERSISTENT ACTOR

- `val receiveRecover: Receive =`
`???`
- `val receiveCommand: Receive =`
`???`

RECEIVECOMMAND

```
case object Snapshot
case class Increase(n: Int)
case class Increased(n: Int)

var state = 0

val receiveCommand: Receive = {
    case Increase(n) =>
        persist(Increased(n)) { event =>
            state = state + n
        }
    case Snapshot => saveSnapshot(state)
}
```


RECEIVERRECOVER

```
val receiveRecover: Receive = {  
  case Increased(n) => state = state + n  
  case SnapshotOffer(_, snapshot: Int) => state = snapshot  
}
```

PERSIST

```
def persist[A](event: A, handler: Procedure[A]): Unit
```

- persists event, then runs the handler
- handler can mutate state
- it is safe to call sender()
- message stashing

EVENT STORAGE IDENTIFIER

- mandatory
- identifies collection used to store events
- override def persistenceld = "never, ever change this value"

AT LEAST ONCE DELIVERY

- trait `AtLeastOnceDelivery` extends `PersistentActor` with `AtLeastOnceDeliveryLike`
- `def deliver(destination: ActorPath)(deliveryIdToMessage: Long ⇒ Any): Unit`
- `def confirmDelivery(deliveryId: Long): Boolean`
- `def redeliverInterval: FiniteDuration`
- `akka.persistence.at-least-once-delivery.redeliver-interval`

AT LEAST ONCE DELIVERY - DELIVERYID

- deliveryId - strictly monotonically increasing sequence number without gaps
- It is not possible to use custom deliveryId
- The trait has a state consisting of unconfirmed messages and a sequence number

AT LEAST ONCE DELIVERY - RECOVERY

- You must store events corresponding to deliver and confirmDelivery calls
- Calls to deliver will not send out messages
- Messages will be sent later if no matching confirmDelivery calls have been performed

SHUTDOWN

- do not use PoisonPill!
- actor would stop before processing stashed messages
- explicit shutdown message
 - case object Stop
 - case object Shutdown
 - case object GoHomeAkkaYouAreDrunk

RELAXED CONSISTENCY REQUIREMENTS

- "(...) process commands as fast as you can, assuming that the Event will eventually be persisted and handled properly in the background, retroactively reacting to persistence failures if needed"
 - to implement the "command sourcing" pattern (store commands instead of events)
- ```
def persistAsync[A](event: A)(handler: Procedure[A]): Unit
```



# PERSISTENT QUERY

- receives updates from an actor
- has own storage - identified by the "viewId" parameter

# SNAPSHOT

- saveSnapshot
- SaveSnapshotSuccess / SaveSnapshotFailure
- final case class SnapshotMetadata(persistenceId: String, sequenceNr: Long, timestamp: Long = 0L)
- SnapshotOffer(metadata, offeredSnapshot)

# SNAPSHOT - RECOVERY

```
override def recovery = Recovery(fromSnapshot = SnapshotSelectionCriteria(
 maxSequenceNr = 457L,
 maxTimestamp = System.currentTimeMillis))
```

---

- SnapshotSelectionCriteria.Latest - default
- SnapshotSelectionCriteria.None - disable snapshot-based recovery
- if no snapshot matched => replay all messages

# RECOVERY

- automatic
- on start and restart
- recovery completed  
message
- stashes incoming  
messages

# EVENT ADAPTERS

- schema evolution (version migration)
- separating domain and data model
- event adapter bindings (conf file)

# TESTING

- `akka.persistence.journal.leveldb.native = off`, or
- `akka.persistence.journal.leveldb-shared.store.native = off`
- Not possible to use `TestActorRef`
- Use `TestKit` instead

# BOOKS AND LINKS

- <http://doc.akka.io/docs/akka/snapshot/scala/persistence.html>
- <http://martinfowler.com/eaDev/EventSourcing.html>
- Vaughn Vernon - Implementing Domain-Driven Design

**WORKSHOP TIME!**